

Team John Hall

Chiew Jun Hao | Ku Ping Cheng | Mu Wenchuan | Zhang Hao

CNN Structure

	Layers				
	1	2	3	4	5
Type	Conv + Pool(max)	Conv + Pool(max)	Conv + Pool(max)	Fully Connected	Fully Connected
Channels	32	32	64	64	2
Filter Size	5 x 5	5 x 5	5 x 5	-	-
Convolution Stride	1 x 1	1 x 1	1 x 1	-	-
Pooling Size	3 x 3	3 x 3	3 x 3	-	-
Pooling Stride	2 x 2	2 x 2	2 x 2	-	-
Padding Size	2 x 2	2 x 2	2 x 2	-	-

Training

The following settings / parameters were used during training:

Fix learning rate: 0.0001

Dropout probability: 0.7

Training batch size: 128

Optimizer: Adam Optimizer

Loss Function: Cross Entropy

Number of patches generated from each image: 100

Patch size: 64

Data Manipulation:

Image is thumbnailed to half the size, because we found that this improves the quality of the image. Besides that, we also subtract the mean of the image RGB value from the image.

We wanted to introduce some noise during training so the network will become more resilient during operation. Below are some of the methods of doing that:

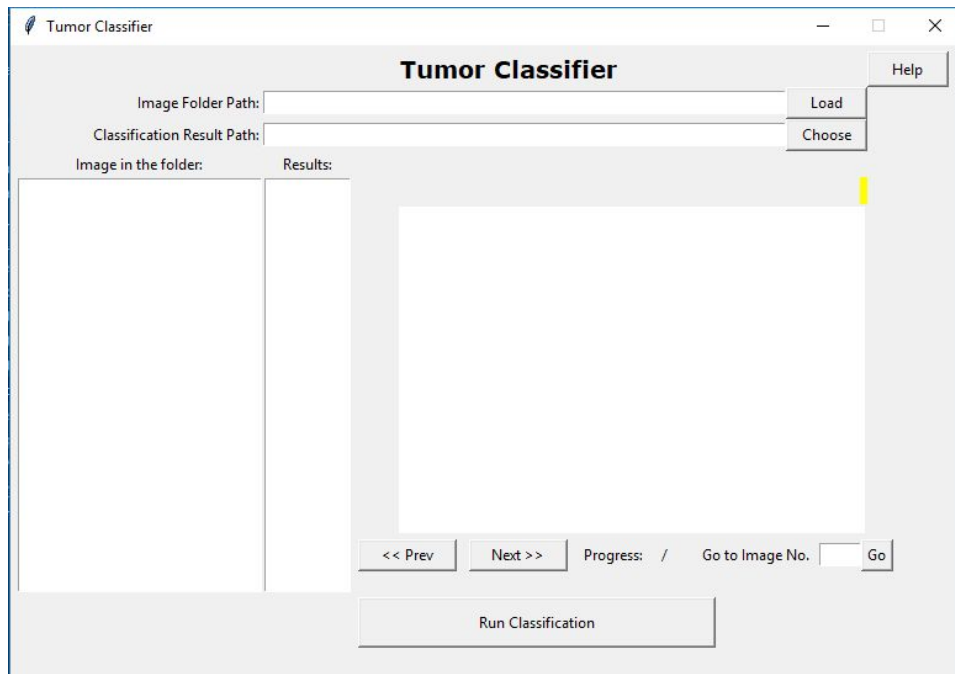
1. Rotate image

Image rotate probability:	Original:	0.4
	90 degrees:	0.2
	180 degrees:	0.2
	270 degrees:	0.2


2. Image is mirrored with a probability of 0.5

GUI

The GUI is made using Tkinter. It takes in a folder path that contains all the images that need to be classified and output path for the classification result.



Once the files are loaded, click Run Classification to classify all images. And the classification result will be shown.



Achieved Accuracy

We chose 100X images during training, validating and testing and achieved the following accuracy:

split	Validation accuracy	Test accuracy
1	80.95%	88.20%
2	68.46%	86.28%
3	63.74%	83.16%
4	80.47%	77.64%
5	93.90%	79.47%
	Average	82.95%